

SpaceX Final Assignment

Bayu Fathurrahman

Executive Summary

- Methodologies
 - Data Collection with API
 - Data Wrangling
 - EDA with SQL
 - EDA and Data Viz with pandas
 - Mapping with Folium
 - Machine Learning Prediction
- Results
 - EDA results
 - Classification modelling results

Table of Contents

- Introduction
- Methodology
- Results
- Discussion
- Conclusion
- Appendix

Introduction

In this capstone, we will predict if the Falcon 9 first stage will land successfully. SpaceX advertises Falcon 9 rocket launches on its website with a cost of 62 million dollars; other providers cost upward of 165 million dollars each, much of the savings is because SpaceX can reuse the first stage. Therefore if we can determine if the first stage will land, we can determine the cost of a launch. This information can be used if an alternate company wants to bid against SpaceX for a rocket launch.

Problems:

- Do we have an increasing rate of success over time?
- How does payload affects the success rate of the landing?
- Which features contribute to the success rate of the landing?

Methodology

Data collection with API

DATA COLLECTION

```
1 spacex_url="https://api.spacexdata.com/v4/launches/past"
2 response = requests.get(spacex_url)
3 # print(response.content)
```

[9]

Python

```
1 static_json_url="https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DS0321EN-SkillsNetwork/datasets/API_call_spacex_api.json"
2 response = requests.get(static_json_url)
3 response.status_code
```

[10]

Python

... 200

```
1 # Use json_normalize meethod to convert the json result into a dataframe
2 data = pd.json_normalize(response.json())
3 data.head()
```

[11]

Python

... static_fire_date_utc static_fire_date_unix tbd net window rocket success details crew ships capsules payloads launchpad auto

| | | | | | | | | | | | | | |
|---|--------------------------|--------------|-------|-------|-----|--------------------------|-------|--|----|----|----|----------------------------|--------------------------|
| 0 | 2006-03-17T00:00:00.000Z | 1.142554e+09 | False | False | 0.0 | 5e9d0d95eda69955f709d1eb | False | Engine failure at 33 seconds and loss of vehicle | [] | [] | [] | [5eb0e4b5b6c3bb0006eeb1e1] | 5e9e4502f5090995de566f86 |
|---|--------------------------|--------------|-------|-------|-----|--------------------------|-------|--|----|----|----|----------------------------|--------------------------|

Successful first stage burn and transition to second

Data Wrangling

```
1 # Lets take a subset of our dataframe keeping only the features we want and the flight number, and date_utc.
2 data = data[['rocket', 'payloads', 'launchpad', 'cores', 'flight_number', 'date_utc']]
3
4 # We will remove rows with multiple cores because those are falcon rockets with 2 extra rocket boosters and rows that have multiple payloads in a single rocket.
5 data = data[data['cores'].map(len)==1]
6 data = data[data['payloads'].map(len)==1]
7
8 # Since payloads and cores are lists of size 1 we will also extract the single value in the list and replace the feature.
9 data['cores'] = data['cores'].map(lambda x : x[0])
10 data['payloads'] = data['payloads'].map(lambda x : x[0])
11
12 # We also want to convert the date_utc to a datetime datatype and then extracting the date leaving the time
13 data['date'] = pd.to_datetime(data['date_utc']).dt.date
14
15 # Using the date we will restrict the dates of the launches
16 data = data[data['date'] <= datetime.date(2020, 11, 13)]
```

Python

```
1 # Dropping Falcon 1 and resetting FlightNumber
2 data_falcon9.drop(data_falcon9[data_falcon9["BoosterVersion"] != "Falcon 9"].index, inplace=True)
3 data_falcon9.loc[:, 'FlightNumber'] = list(range(1, data_falcon9.shape[0]+1))
4 data_falcon9.reset_index(drop=True, inplace=True)
5
6 # Replace the np.nan values with its mean value
7 data_falcon9["PayloadMass"].replace(np.nan, data_falcon9["PayloadMass"].mean(), inplace=True)
```

Python

```
1 data_falcon9["Class"] = data_falcon9["Outcome"].apply(lambda x: 1 if x in ["True ASDS", "True RTLS", "True Ocean"] else 0)
```

Python

```
1 data_falcon9["Class"].value_counts(True)
```

Python

```
1    0.666667
0    0.333333
Name: Class, dtype: float64
```

EDA with SQL

```
1 %%sql
2
3 select distinct(launch_site) from spacextbl;
```

✓ 0.5s

* sqlite:///my_data2.db

Done.

Launch_Site

CCAFS LC-40

VAFB SLC-4E

KSC LC-39A

CCAFS SLC-40

5 records where launch sites begin with the string 'CCA'

```
1 %%sql
2
3 select * from spacextbl
4 where launch_site like "CCA%"
5 limit 5;
```

✓ 0.4s

* sqlite:///my_data2.db

Done.

| Date | Time (UTC) | Booster_Version | Launch_Site | Payload | PAYLOAD_MASS_KG_ | Orbit | Customer | Mission_Outcome | Landing_Outcome |
|------------|------------|-----------------|-------------|---|------------------|-----------|-----------------|-----------------|---------------------|
| 04-06-2010 | 18:45:00 | F9 v1.0 B0003 | CCAFS LC-40 | Dragon Spacecraft Qualification Unit | 0 | LEO | SpaceX | Success | Failure (parachute) |
| 08-12-2010 | 15:43:00 | F9 v1.0 B0004 | CCAFS LC-40 | Dragon demo flight C1, two CubeSats, barrel of Brouere cheese | 0 | LEO (ISS) | NASA (COTS) NRO | Success | Failure (parachute) |
| 22-05-2012 | 07:44:00 | F9 v1.0 B0005 | CCAFS LC-40 | Dragon demo flight C2 | 525 | LEO (ISS) | NASA (COTS) | Success | No attempt |
| 08-10-2012 | 00:35:00 | F9 v1.0 B0006 | CCAFS LC-40 | SpaceX CRS-1 | 500 | LEO (ISS) | NASA (CRS) | Success | No attempt |
| 01-03-2013 | 15:10:00 | F9 v1.0 B0007 | CCAFS LC-40 | SpaceX CRS-2 | 677 | LEO (ISS) | NASA (CRS) | Success | No attempt |

Python

total payload mass carried by boosters launched by NASA (CRS)

```
1 %%sql
2
3 select customer, sum(payload_mass_kg_) as "total payload" from spacextbl
4 where customer = "NASA (CRS)";
```

✓ 0.4s

* sqlite:///my_data2.db

Done.

Customer total payload

NASA (CRS) 45596

average payload mass carried by booster version F9 v1.1

```
1 %%sql
2
3 select booster_version, avg(payload_mass_kg_) as "average payload" from spacextbl
4 where booster_version = "F9 v1.1";
```

✓ 0.4s

* sqlite:///my_data2.db

Done.

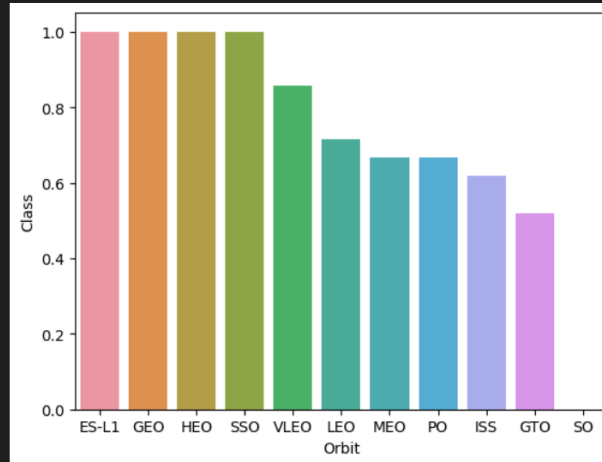
Booster_Version average payload

F9 v1.1 2928.4

EDA with pandas

```
1 df_orbit = df.groupby("Orbit").mean()[["Class"]].sort_values("Class", ascending=False)
2 sns.barplot(x=df_orbit.index, y=df_orbit["Class"]);
```

Python



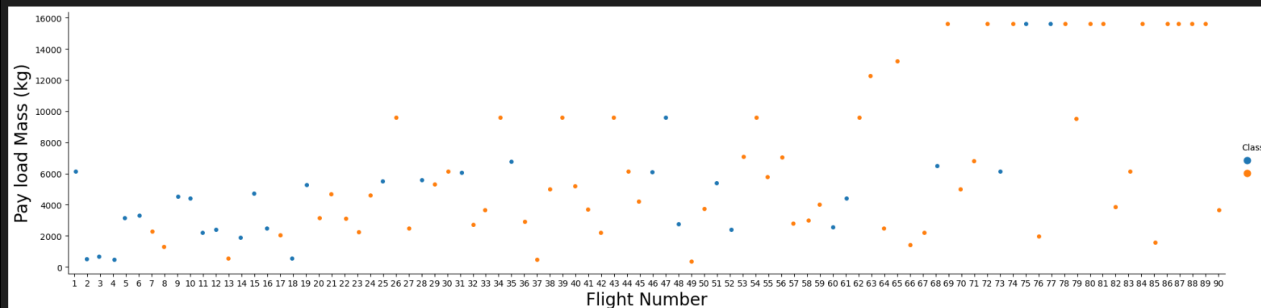
```
1 df[df["Orbit"] == "so"]
```

Python

| FlightNumber | | Date | BoosterVersion | PayloadMass | Orbit | LaunchSite | Outcome | Flights | GridFins | Reused | Legs | LandingPad | Block | ReusedCount | Serial | Longitude | Latitude | Class |
|--------------|----|------------|----------------|-------------|-------|------------|--------------|---------|----------|--------|-------|------------|-------|-------------|--------|------------|-----------|-------|
| 72 | 73 | 2020-01-19 | Falcon 9 | 6123.547647 | SO | KSC LC 39A | None None | 4 | False | True | False | NaN | 5 | 3 | B1046 | -80.603956 | 28.608058 | 0 |

SO orbit has a success rate of 0% because there's only 1 launch attempt and it hasn't succeed. Meanwhile, Launch on orbit ES-L1, GEO, HEO, and SSO have 100% success rate

```
1 style.use("default")
2 sns.catplot(y="PayloadMass", x="flightNumber", hue="Class", data=df, aspect=4)
3 plt.xlabel("Flight Number", fontsize=20)
4 plt.ylabel("Pay load Mass (kg)", fontsize=20)
5 plt.show()
```

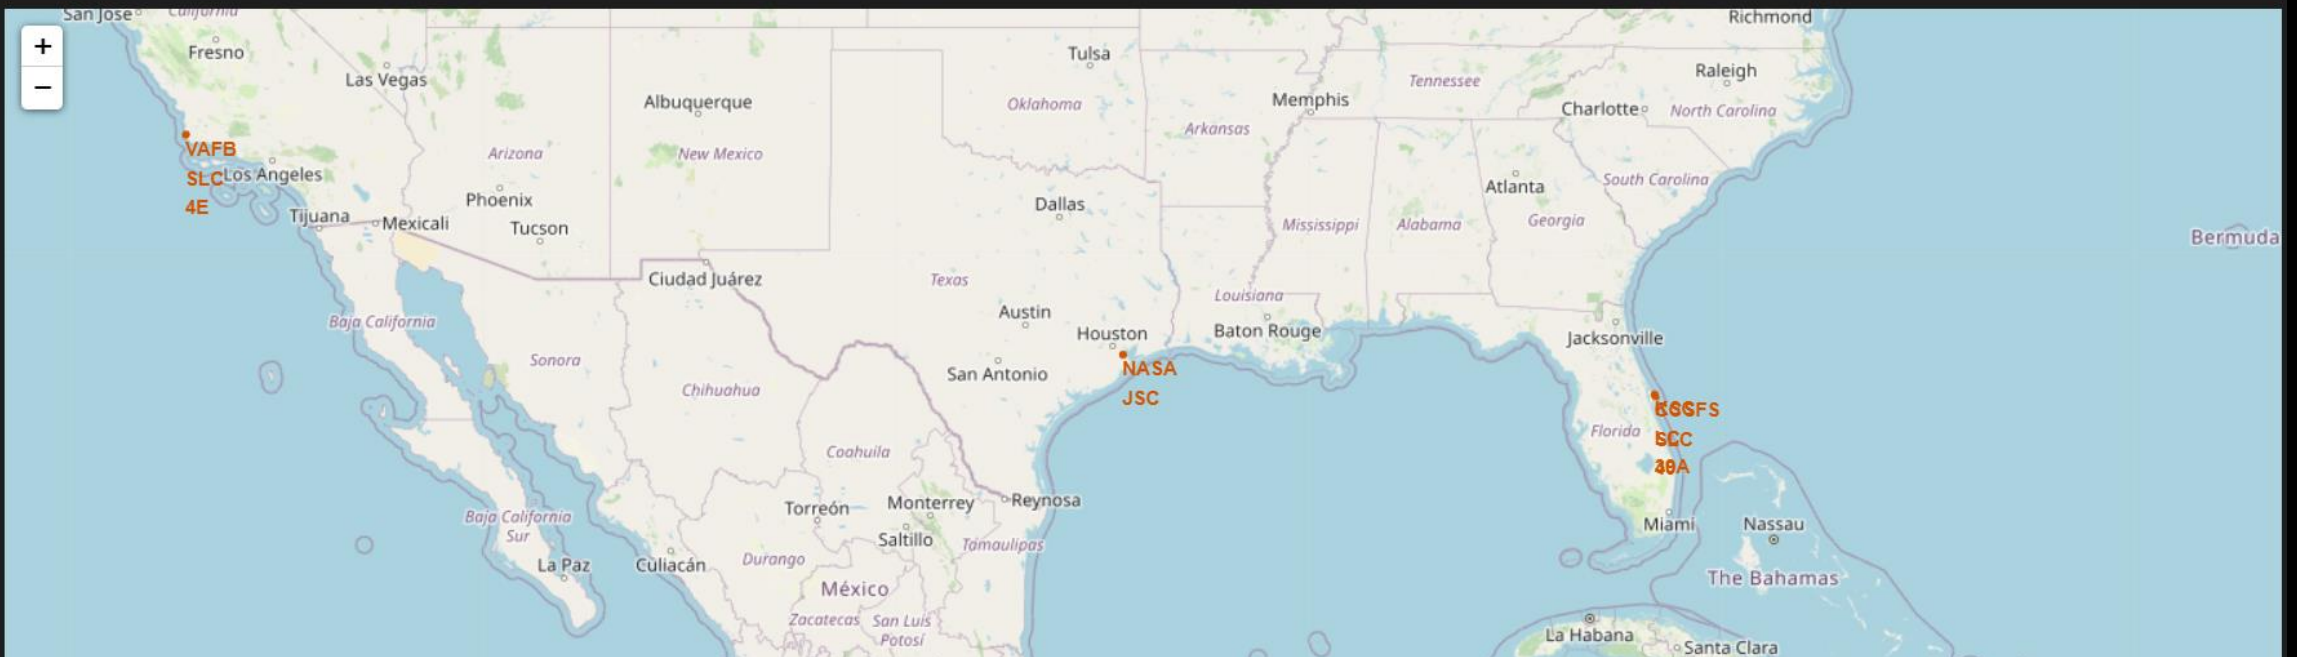


We see that as the flight number increases, the first stage is more likely to land successfully. The payload mass is also important; it seems the more massive the payload, the less likely the first stage will return.

Mapping with Folium

```
1 nasa_coordinate = [28.561857, -80.577366]
2 # Create a circle at NASA Johnson Space Center's coordinate with a popup label showing its name
3 circle = folium.Circle(nasa_coordinate, radius=1000, color='#d35400', fill=True).add_child(folium.Popup('CCSFS SLC 40'))
4 # Create a circle at NASA Johnson Space Center's coordinate with a icon showing its name
5 marker = folium.map.Marker(
6     nasa_coordinate,
7     # Create an icon as a text label
8     icon=DivIcon(
9         icon_size=(20,20),
10        icon_anchor=(0,0),
11        html='<div style="font-size: 12; color:#d35400;"><b>%s</b></div>' % 'CCSFS SLC 40',
12    )
13 )
14 site_map.add_child(circle)
15 site_map.add_child(marker)
```

Python



Modelling

```
1 from sklearn import preprocessing
2 transform = preprocessing.StandardScaler()
```

```
1 transform.fit_transform(X)
```

```
array([[ -1.71291154,  0.          , -0.65391284, ..., -0.21566555,
        -0.18569534, -0.10599979],
       [ -1.67441914, -1.18972425, -0.65391284, ..., -0.21566555,
        -0.18569534, -0.10599979],
       [ -1.63592675, -1.15742336, -0.65391284, ..., -0.21566555,
        -0.18569534, -0.10599979],
       ...,
       [  1.63592675,  2.01380177,  3.49060516, ..., -0.21566555,
        -0.18569534, -0.10599979],
       [  1.67441914,  2.01380177,  1.00389436, ..., -0.21566555,
         5.38516481, -0.10599979],
       [  1.71291154, -0.51905572, -0.65391284, ..., -0.21566555,
        -0.18569534,  9.43398113]])
```

```
1 X_train, X_test, y_train, y_test = train_test_split(
2     X,
3     y,
4     stratify=y,
5     test_size = 0.2,
6     random_state = 2)
```

LogReg

```
1 parameters = {"C": [0.01, 0.1, 1], 'penalty': ['l2'], 'solver': ['lbfgs']} # l1 lasso l2 ridge
2 lr = LogisticRegression()
3
4 logreg_cv = GridSearchCV(
5     estimator = lr,
6     param_grid = parameters,
7     cv = 10,
8     scoring = "accuracy",
9     return_train_score = True,
10    n_jobs = -1
11 )
12
13 logreg_cv.fit(X_train, y_train)
14
15 print("tuned hpyerparameters :(best parameters) ", logreg_cv.best_params_)
16 print("accuracy :", logreg_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters) {'C': 0.1, 'penalty': 'l2', 'solver': 'lbfgs'}
accuracy : 0.8214285714285714
```

```
1 logreg_cv_best = logreg_cv.best_estimator_
2 logreg_cv_best.fit(X_train, y_train)
3
4 y_test_pred = logreg_cv_best.predict(X_test)
5 style.use("default")
6 ConfusionMatrixDisplay(confusion_matrix(y_test, y_test_pred)).plot();
```

KNN

```
1 parameters = {'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
2               'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
3               'p': [1, 2]}
4 KNN = KNeighborsClassifier(n_jobs=-1)
5
6 knn_cv = GridSearchCV(
7     estimator = KNN,
8     param_grid = parameters,
9     cv = 10,
10    scoring = "accuracy",
11    return_train_score = True,
12    n_jobs = -1
13 )
14
15 knn_cv.fit(X_train, y_train)
16
17 print("tuned hpyerparameters :(best parameters) ", knn_cv.best_params_)
18 print("accuracy :", knn_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters) {'algorithm': 'auto', 'n_neighbors': 1, 'p': 2}
accuracy : 0.6535714285714286
```

```
1 knn_cv_best = knn_cv.best_estimator_
2 knn_cv_best.fit(X_train, y_train)
3
4 y_test_pred = knn_cv_best.predict(X_test)
5 style.use("default")
6 ConfusionMatrixDisplay(confusion_matrix(y_test, y_test_pred)).plot();
```

Decision Tree

```
1 parameters = {'criterion': ['gini', 'entropy'],
2               'splitter': ['best', 'random'],
3               'max_depth': [2*n for n in range(1, 10)],
4               'max_features': ['auto', 'sqrt'],
5               'min_samples_leaf': [1, 2, 4],
6               'min_samples_split': [2, 5, 10]}
7 tree = DecisionTreeClassifier(random_state=42)
```

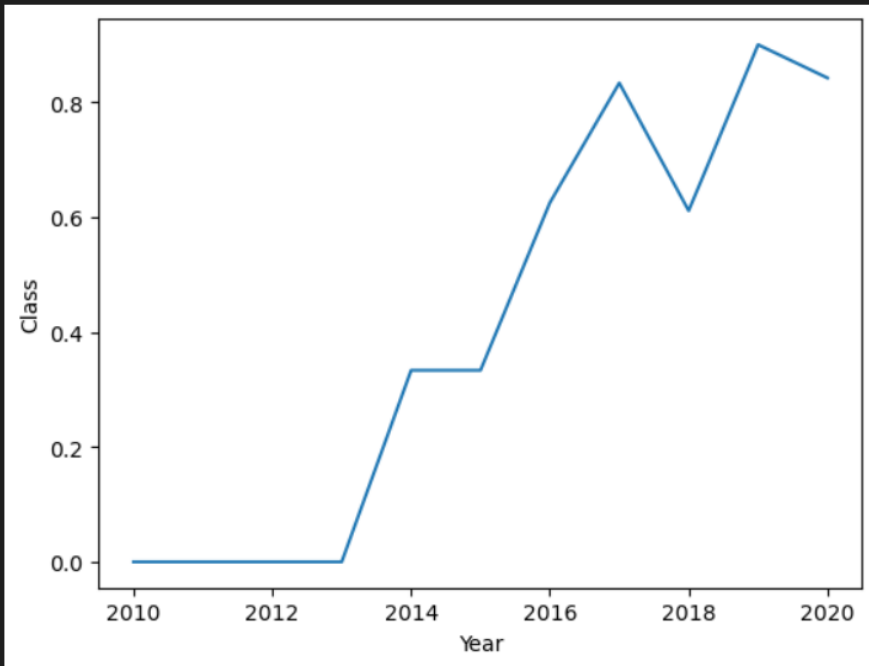
```
8
9 tree_cv = GridSearchCV(
10     estimator = tree,
11     param_grid = parameters,
12     cv = 10,
13     scoring = "accuracy",
14     return_train_score = True,
15     n_jobs = -1
16 )
17
18 tree_cv.fit(X_train, y_train)
19
20 print("tuned hpyerparameters :(best parameters) ", tree_cv.best_params_)
21 print("accuracy :", tree_cv.best_score_)
```

```
tuned hpyerparameters :(best parameters) {'criterion': 'gini', 'max_depth': 8,
accuracy : 0.8785714285714287
```

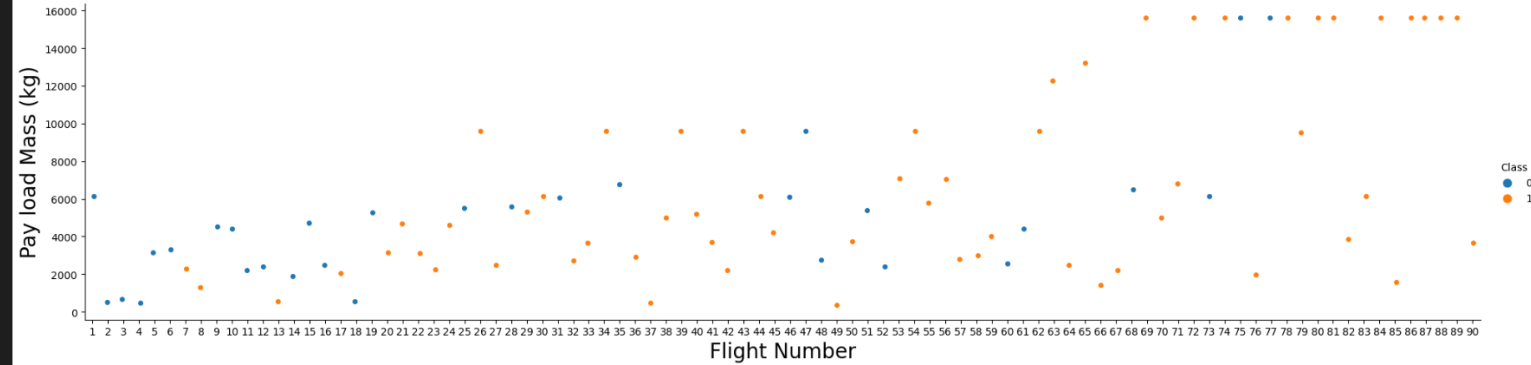
```
1 tree_cv_best = tree_cv.best_estimator_
2 tree_cv_best.fit(X_train, y_train)
3
4 y_test_pred = tree_cv_best.predict(X_test)
5 style.use("default")
6 ConfusionMatrixDisplay(confusion_matrix(y_test, y_test_pred)).plot();
```

Results

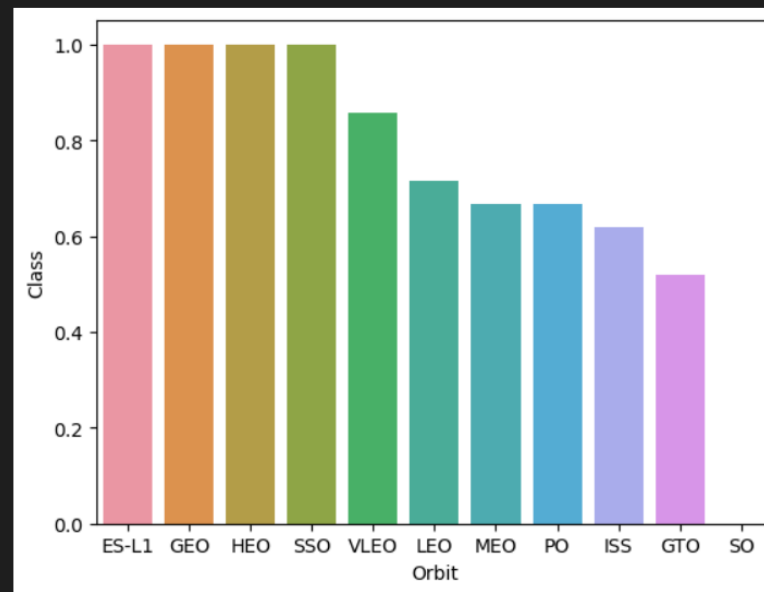
EDA Results



success rate since 2013 kept increasing until 2020



We see that as the flight number increases, the first stage is more likely to land successfully. The payload mass is also important; it seems the more massive the payload, the less likely the first stage will return.



SO orbit has a success rate of 0% because there's only 1 launch attempt and it hasn't succeed. Meanwhile, Launch on orbit ES-L1, GEO, HEO, and SSO have 100% success rate

EDA Results

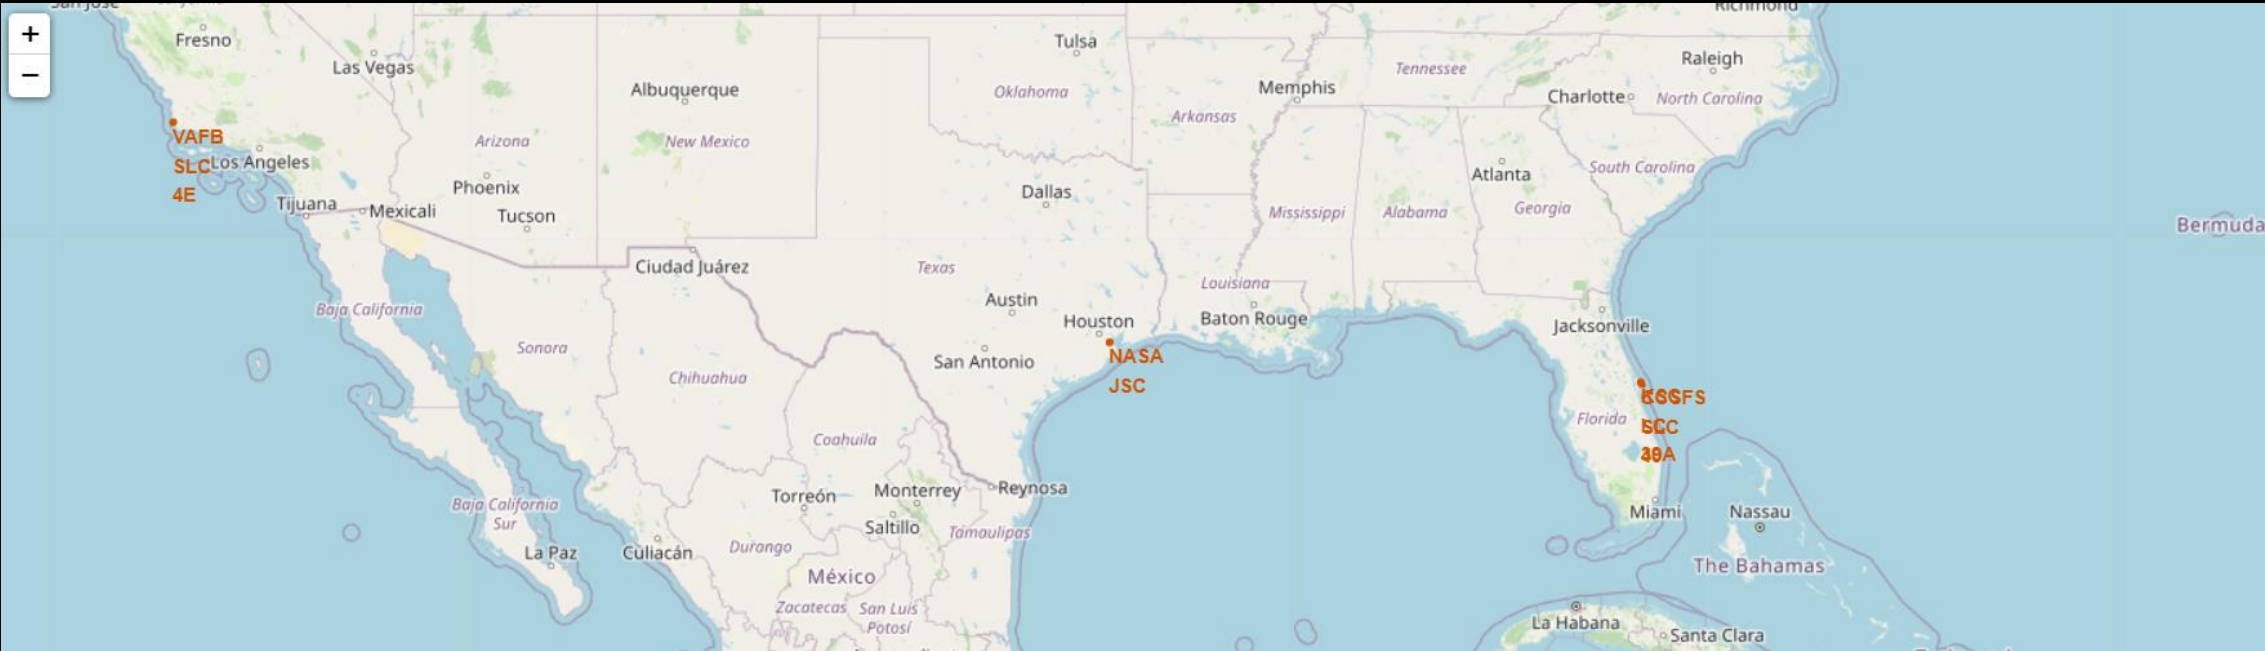
Name of the launch sites

| Launch_Site |
|--------------|
| CCAFS LC-40 |
| VAFB SLC-4E |
| KSC LC-39A |
| CCAFS SLC-40 |

Total Payload Carried by NASA

| Customer | total payload |
|------------|---------------|
| NASA (CRS) | 45596 |

Launch site locations



Modelling Results

Logreg Test Accuracy (Best Test Result)

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.50 | 0.67 | 6 |
| 1 | 0.80 | 1.00 | 0.89 | 12 |
| accuracy | | | 0.83 | 18 |
| macro avg | 0.90 | 0.75 | 0.78 | 18 |
| weighted avg | 0.87 | 0.83 | 0.81 | 18 |

Logreg Train Accuracy (Best fit)

accuracy : 0.8214285714285714

Decision Tree Test Accuracy (Overfitting)

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.33 | 0.33 | 0.33 | 6 |
| 1 | 0.67 | 0.67 | 0.67 | 12 |
| accuracy | | | 0.56 | 18 |
| macro avg | 0.50 | 0.50 | 0.50 | 18 |
| weighted avg | 0.56 | 0.56 | 0.56 | 18 |

Decision Tree Train Accuracy
(best TRAIN result)

accuracy : 0.8785714285714287

KNN Test Accuracy (Underfitting)

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.38 | 0.50 | 0.43 | 6 |
| 1 | 0.70 | 0.58 | 0.64 | 12 |
| accuracy | | | 0.56 | 18 |
| macro avg | 0.54 | 0.54 | 0.53 | 18 |
| weighted avg | 0.59 | 0.56 | 0.57 | 18 |

KNN Train Accuracy
(n_neighbors=1 from gridsearch)

accuracy : 0.6535714285714286

NO RESULT FOR SVM MODEL BECAUSE IT CRASHED MY PC AFTER TRYING TO RUN IT FOR 30 MINS

code for svm written below:

Conclusion

- Success rate keeps increasing over time, as more and more launch is being done.
- Heavier payloads seems to have a better success on landing, however this might happen because of the fact that most of the heavy payloads are launched only after flight number is relatively high (correlates with our 1st point in this conclusion).
- Launch on orbit ES-L1, GEO, HEO, and SSO have 100% success rate.
- We find that Logistic Regression model is best fit for this dataset. (83% accuracy)

Thank
you

